Package 'isodisregSD'

November 14, 2023

Title IDR for distributional input and iso-based CRPS decomposition

Version 1.0

Description Computes IDR for different types of distributional input data. Yields calibrated forecasts which are used to compute iso-based CRPS decomposition.

License MIT + file LICENSE

Encoding UTF-8

Roxygen list(markdown = TRUE)

RoxygenNote 7.2.3

LinkingTo Rcpp

Imports Rcpp

URL https://github.com/evwalz/isodisregSD

BugReports https://github.com/evwalz/isodisregSD/issues

Suggests knitr,
rmarkdown

VignetteBuilder knitr

R topics documented:

isodisregSD-package *Isotonic distributional Regression (IDR) for distributional data*

Description

Isotonic distributional Regression (IDR) for distributional data is a nonparametric method to estimate conditional distributions under stochastic order constraints.

How does it work?

Link to Preprint on ArXiv as soon as available

bscore *Brier score for forecast probability of threshold exceedance*

Description

Computes the Brier score of forecast probabilities for exceeding given thresholds.

Usage

```
bscore(predictions, thresholds, y)
```

Arguments

predictions	either an object of class idrsd (output of predict.idrcal), or a data.frame of numeric variables. In the latter case, the CDF is computed using the empirical distribution of the variables in predictions.
thresholds	numeric vector of thresholds at which the CDF will be evaluated.
y	a numeric vector of obervations of the same length as the number of predictions, or of length 1. In the latter case, y will be used for all predictions.

Details

The Brier score for the event of exceeding a given threshold z is defined as

$$(1\{y > z\} - P(y > z))^2$$

where *y* is the observation and *P(y > z)* the forecast probability for exceeding the threshold z.

Value

A matrix of the Brier scores for the desired thresholds, one column per threshold.

References

Gneiting, T. and Raftery, A. E. (2007), 'Strictly proper scoring rules, prediction, and estimation', Journal of the American Statistical Association 102(477), 359-378

See Also

predict.idrcal, cdf

Examples

```
data("rain")

## IDRsd based on ensemble forecast

ensemble <- rain[1:(3 * 365), 3:54, drop = FALSE]
y <- rain[1:(3 * 365), "obs"]

fit <- idrsd(y = y, X = ensemble, type = 'ensemble')

## Compute Brier score for probability of precipitation
## forecast using data of the next year (out-of-sample predictions)

ytest = rain[(3 * 365 + 1):(4 * 365), "obs"]
ensemble_test = rain[(3 * 365 + 1):(4 * 365), 3:54,drop = FALSE]
predictions <- predict(fit, data = ensemble_test)
score <- bscore(predictions, thresholds = 0, y = ytest)

mean(score)
```

---

cdf *Cumulative distribution function (CDF) of IDRsd or raw forecasts*

Description

Evaluate the the cumulative distribution function (CDF) of IDRsd predictions or of unprocessed forecasts in a data.frame.

Usage

```
cdf(predictions, thresholds)

## S3 method for class 'idrsd'
cdf(predictions, thresholds)

## S3 method for class 'data.frame'
cdf(predictions, thresholds)
```

Arguments

predictions	either an object of class idrsd (output of predict.idrcal), or a data.frame of numeric variables. In the latter case, the CDF is computed using the empirical distribution of the variables in predictions.
thresholds	numeric vector of thresholds at which the CDF will be evaluated.

Details

The CDFs are considered as piecewise constant stepfunctions: If x are the points where the IDR fitted CDF (or the empirical distribution of the forecasts) has jumps and p the corresponding CDF values, then for x[i] <= x < x[i + 1], the CDF at x is p[i].

Value

A matrix of probabilities giving the evaluated CDFs at the given thresholds, one column for each threshold.

See Also

predict.idrcal qpred, bscore

Examples

```
## Data from IDR package:
library(isodistrreg)
data("rain")

## IDRsd based on ensemble forecast

ensemble <- rain[1:(3 * 365), 3:54, drop = FALSE]
y <- rain[1:(3 * 365), "obs"]

fit <- idrsd(y = y, X = ensemble, type = 'ensemble')
predictions <- predict(fit)

## Compute probability of precipitation
1 - cdf(predictions, thresholds = 0)
```

---

crps *Continuous ranked probability score (CRPS)*

Description

Computes the CRPS of IDRsd or raw forecasts.

Usage

```
crps(predictions, y)

## S3 method for class 'idrsd'
crps(predictions, y)

## S3 method for class 'idr'
crps(predictions, y)

## S3 method for class 'data.frame'
crps(predictions, y)
```

Arguments

predictions	either an object of class idrsd (output of predict.idrcal), or a data.frame of numeric variables. In the latter case, the CRPS is computed using the empirical distribution of the variables in predictions.
y	a numeric vector of obervations of the same length as the number of predictions, or of length 1. In the latter case, y will be used for all predictions.

Details

This function uses adapted code taken from the function `crps_edf` of the **scoringRules** package.

Value

A vector of CRPS values.

References

Jordan A., Krueger F., Lerch S. (2018). "Evaluating Probabilistic Forecasts with scoringRules." Journal of Statistical Software. Forthcoming.

Gneiting, T. and Raftery, A. E. (2007), 'Strictly proper scoring rules, prediction, and estimation', Journal of the American Statistical Association 102(477), 359-378

See Also

`predict.idrcal`

Examples

```
data("rain")

## IDRsd based on ensemble forecast

ensemble <- rain[1:(3 * 365), 3:54, drop = FALSE]
y <- rain[1:(3 * 365), "obs"]

fit <- idrsd(y = y, X = ensemble, type = 'ensemble')

## Compute CRPS of forecast using data of the next year
## (out-of-sample predictions)

ytest = rain[(3 * 365 + 1):(4 * 365), "obs"]
ensemble_test = rain[(3 * 365 + 1):(4 * 365), 3:54,drop = FALSE]
predictions <- predict(fit, data = ensemble_test)
idrCrps <- crps(predictions, y = ytest)

## Compare this to CRPS of the raw ensemble of all forecasts (high resolution,
## control and 50 perturbed ensemble forecasts)

rawData <- rain[(3 * 365 + 1):(4 * 365), c("HRES", "CTR", paste0("P", 1:50))]
rawCrps <- crps(rawData, y = ytest)

c("idrsd" = mean(idrCrps), "raw_all" = mean(rawCrps))
```

---

idrsd *Fit IDRsd to training data*

Description

Fits isotonic distributional regression (IDR) to a training dataset with respect to the stochastic order (SD) order.

Usage

```
idrsd(y, X, pars = osqpSettings(verbose = FALSE, eps_abs =
  1e-5, eps_rel = 1e-5, max_iter = 10000L), progress = TRUE)
```

Arguments

y	numeric vector (the response variable).
X	object that fits to specification in type. For type = 'idr', X must be an IDR object, for type = 'ensemble', X must be a data.frame where columns correspond to ensemble members, for type = 'ecdf', X must be a (n x m) matrix with n = length(y) and m = length(grid), where each row corresponds to ECDF values evaluated at grid. For type = 'dis' X is empty, for type = 'normal' and type = 'normal_ab' X is a matrix with 2 columns which represent mu and sigma parameters of the normal distribution.
grid	if type == 'ecdf', than grid is vector of threshold values corresponding to ECDF-values in X
dis_func	if type == 'dis', then a cumulative distribution function must be specified along with its distributional parameters.
type	default is 'ensemble'. Other possibilities are 'idr', 'ecdf', 'dis', 'normal', 'normal_ab'
pars	parameters for quadratic programming optimization (only relevant if X has more than one column), set using osqpSettings.
progress	display progressbar (TRUE, FALSE or 1, 0)?

Details

This function computes the isotonic distributional regression (IDR) of a response y on distributional data X using stochastic order. IDRsd estimates the cumulative distribution function (CDF) of y conditional on X under the assumption of isotonicity. The odering on X is computed using the stochastic order. The conditional CDFs are estimated at each threshold in unique(y). This is the set where the CDFs may have jumps.

Value

An object of class "idrcal" containing the following components:

X	data frame of all distinct values (note the different order!).
y	list of all observed responses in the training data for given X.
cdf	matrix containing the estimated CDFs, one CDF per row, evaluated at thresholds (see next point). The CDF in the ith row corredponds to the estimated conditional distribution of the response given X[i,].
thresholds	the thresholds at which the CDFs in cdf are evaluated. The entries in cdf[,j] are the conditional CDFs evaluated at thresholds[j].
type	type of distributional input data as specified by user
diagnostic	list giving a bound on the precision of the CDF estimation (the maximal downwards-step in the CDF that has been detected) and the fraction of CDF estimations that were stopped at the iteration limit max_iter. Decrease the parameters eps_abs and/or eps_rel or increase max_iter in pars to improve the precision. See osqpSettings for more optimization parameters.

indices	the row indices of X in the original training dataset (used for in-sample predictions with predict.idrfit).
constraints	(in multivariate IDR, NULL otherwise) matrices giving the order constraints for optimization. Used in predict.idrfit.

Note

The function idrsd is only intended for fitting IDR model for a training dataset and storing the results for further processing, but not for prediction or evaluation, which is done using the output of predict.idrcal.

See Also

The S3 method predict.idrcal for predictions based on an IDR fit.

---

isodeco_crps *CRPS decomposition*

Description

Computes the individual components of the iso-based CRPS Decomposition: MSC, DSC and UNC

Usage

```
isodeco_crps(
  y,
  X = NULL,
  grid = NULL,
  dis_func = NULL,
  type = "ensemble",
  inta = NULL,
  intb = NULL,
  pars = osqpSettings(verbose = FALSE, eps_abs = 1e-05, eps_rel = 1e-05, max_iter =
    10000L),
  progress = TRUE,
  ...
)
```

Arguments

y	numeric vector (the response variable).
X	object that fits to specification in type. For type = 'idr', X must be an IDR object, for type = 'ensemble', X must be a data.frame where columns correspond to ensemble members, for type = 'ecdf', X must be a (n x m) matrix with n = length(y) and m = length(grid), where each row corresponds to ECDF values evaluated at grid. For type = 'dis' X is empty, for type = 'normal' and type = 'normal_ab' X is a matrix with 2 columns which represent mu and sigma parameters of the normal distribution.
grid	if type == 'ecdf', than grid is vector of threshold values corresponding to ECDF-values in X

dis_func	if type == 'dis', then a cumulative distribution function must be specified along with its distributional parameters.
type	default is 'ensemble'. Other possibilities are 'idr', 'ecdf', 'dis', 'normal', 'normal_ab'
pars	parameters for quadratic programming optimization (only relevant if X has more than one column), set using osqpSettings.
progress	display progressbar (TRUE, FALSE or 1, 0)?

Details

This function computed the CRPS decomposition of a response vector *y* and a distributional forecast *X*, which can be an ensemble and empirical cumulative distributional function, a normal distribution or any other other closed form CDF specified by its parameters.

Value

A list of CRPS decomposition components: miscalibration (MSC), discrimination (DSC), uncertainty (UNC) and the original CRPS.

Examples

```
#' data("rain")

## IDRsd based on ensemble forecast

ensemble <- rain[1:(3 * 365), 3:54, drop = FALSE]
y <- rain[1:(3 * 365), "obs"]

## Compute CRPS decomposition components
crps_deco <- isodeco_crps(y = y, X = ensemble, type = 'ensemble')
print(crps_deco)
```

---

pit *Probability integral transform (PIT)*

Description

Computes the probability integral transform (PIT) of IDRsd or raw forecasts.

Usage

```
pit(predictions, y, randomize = TRUE, seed = NULL)

## S3 method for class 'idrsd'
pit(predictions, y, randomize = TRUE, seed = NULL)

## S3 method for class 'data.frame'
pit(predictions, y, randomize = TRUE, seed = NULL)
```

Arguments

predictions	either an object of class idrsd (output of predict.idrcal), or a data.frame of numeric variables. In the latter case, the PIT is computed using the empirical distribution of the variables in predictions.
y	a numeric vector of obervations of the same length as the number of predictions.
randomize	PIT values should be randomized at discontinuity points of the predictive CDF (e.g. at zero for precipitation forecasts). Set randomize = TRUE to randomize.
seed	argument to set.seed for random number generation (if randomize is TRUE).

Value

Vector of PIT values.

References

Gneiting, T., Balabdaoui, F. and Raftery, A. E. (2007), 'Probabilistic forecasts, calibration and sharpness', Journal of the Royal Statistical Society: Series B (Statistical Methodology) 69(2), 243-268.

See Also

predict.idrcal

Examples

```
data("rain")
require("graphics")


ensemble <- rain[1:(3 * 365), 3:54, drop = FALSE]
y <- rain[1:(3 * 365), "obs"]

fit <- idrsd(y = y, X = ensemble, type = 'ensemble')

## Assess forecast using data of next 2
## years and compare to calibration of the raw ensemble

ytest = rain[(3 * 365 + 1):(5 * 365), "obs"]
ensemble_test = rain[(3 * 365 + 1):(5 * 365), 3:54,drop = FALSE]
predictions <- predict(fit, data = ensemble_test)
idrPit <- pit(predictions, ytest, seed = 123)

rawData <- rain[(3 * 365 + 1):(5 * 365), c("HRES", "CTR", paste0("P", 1:50))]
rawPit <- pit(rawData, ytest, seed = 123)

hist(idrPit, xlab = "Probability Integral Transform",
  ylab = "Density", freq = FALSE, main = "Calibrated ensemble")
hist(rawPit, xlab = "Probability Integral Transform",
  ylab = "Density", freq = FALSE, main = "Raw ensemble")
```

`plot.idrsd` *Plot IDRsd predictions*

Description

Plot an IDRsd predictive CDF.

Usage

```
## S3 method for class 'idrsd'
plot(
  x,
  index = 1,
  bounds = TRUE,
  col.cdf = "black",
  col.bounds = "blue",
  lty.cdf = 1,
  lty.bounds = 3,
  xlab = "Threshold",
  ylab = "CDF",
  main = "IDR predictive CDF",
  ...
)
```

Arguments

`x`	object of class `idrsd` (output of `predict.idrcal`).
`index`	index of the prediction in `x` for which a plot is desired.
`bounds`	whether the bounds should be plotted or not (see `predict.idrfit` for details about the meaning of the bounds).
`col.cdf`	color of the predictive CDF.
`col.bounds`	color of the bounds.
`lty.cdf`	linetype of the predictive CDF.
`lty.bounds`	linetype of the CDF bounds.
`xlab`	label for x axis.
`ylab`	label for y axis.
`main`	main title.
`...`	further arguments to `plot.stepfun` or `plot`.

Value

The data based on which the plot is drawn (returned invisible).

See Also

`predict.idrcal`

Examples

```
data("rain")
require("graphics")


ensemble <- rain[1:(3 * 365), 3:54, drop = FALSE]
y <- rain[1:(3 * 365), "obs"]

## Fit IDRsd and plot the predictive CDF

fit <- idrsd(y = y, X = ensemble, type = 'ensemble')
ensemble_test = rain[(3 * 365 + 1), 3:54,drop = FALSE]
pred <- predict(fit, data = ensemble_test)
plot(pred)
```

`predict.idrcal` *Predict method for IDRsd fits*

Description

Prediction based on IDRsd model fit.

Usage

```
## S3 method for class 'idrcal'
predict(
  object,
  data = NULL,
  grid = NULL,
  digits = 3,
  interpolation = "linear",
  ...
)
```

Arguments

`object`	IDRsd fit (object of class `"idrcal"`).
`data`	optional object with which to predict. In-sample predictions are returned if this is omitted.
`digits`	number of decimal places for the predictive CDF.
`...`	included for generic function consistency.

Details

If the variables `data` for which predictions are desired are already contained in the training dataset `X` for the fit, `predict.idrfit` returns the corresponding in-sample prediction. Otherwise monotonicity is used to derive upper and lower bounds for the predictive CDF, and the predictive CDF is a pointwise average of these bounds.

Value

A list of predictions. Each prediction is a `data.frame` containing the following variables:

`points`	the points where the predictive CDF has jumps.
`cdf`	the estimated CDF evaluated at the `points`.
`lower, upper`	(only for out-of-sample predictions) bounds for the estimated CDF, see 'Details' above.

The output has the attribute `incomparables`, which gives the indices of all predictions for which the climatological forecast is returned because the forecast variables are not comparable to the training data.

See Also

`idrsd` to fit IDRsd to training data.

`cdf`, `qpred` to evaluate the CDF or quantile function of IDR predictions.

`bscore`, `qscore`, `crps`, `pit` to compute Brier scores, quantile scores, the CRPS and the PIT of IDR predictions.

`plot` to plot IDR predictive CDFs.

---

qpred *Quantile function of IDRsd or raw forecasts*

Description

Evaluate the the quantile function of IDRsd predictions or of unprocessed forecasts in a `data.frame`.

Usage

```
qpred(predictions, quantiles)

## S3 method for class 'idrsd'
qpred(predictions, quantiles)

## S3 method for class 'data.frame'
qpred(predictions, quantiles)
```

Arguments

`predictions`	either an object of class `idrsd` (output of `predict.idrcal`), or a `data.frame` of numeric variables. In the latter case, quantiles are computed using the empirical distribution of the variables in `predictions`.
`quantiles`	numeric vector of desired quantiles.

Details

The quantiles are defined as lower quantiles, that is,

$$q(u) = inf(x : cdf(x) >= u).$$

Value

A matrix of forecasts for the desired quantiles, one column per quantile.

See Also

predict.idrcal, cdf, qscore

Examples

```
## Data from IDR package:
library(isodistrreg)
data("rain")

## IDRsd based on ensemble forecast

ensemble <- rain[1:(3 * 365), 3:54, drop = FALSE]
y <- rain[1:(3 * 365), "obs"]

fit <- idrsd(y = y, X = ensemble, type = 'ensemble')
predictions <- predict(fit)

## Compute 95%-quantile forecast
qpred(predictions, quantiles = 0.95)
```

---

qscore *Quantile scores for IDRsd or raw forecasts*

Description

Computes quantile scores of IDRsd quantile predictions or of quantile predictions from raw forecasts in a data.frame.

Usage

```
qscore(predictions, quantiles, y)
```

Arguments

predictions	either an object of class idrsd (output of predict.idrcal), or a data.frame of numeric variables. In the latter case, quantiles are computed using the empirical distribution of the variables in predictions.
quantiles	numeric vector of desired quantiles.
y	a numeric vector of obervations of the same length as the number of predictions, or of length 1. In the latter case, y will be used for all predictions.

Details

The quantile score of a forecast x for the u-quantile is defined as

$$2(1x > y - u)(x - y),$$

where y is the observation. For $u = 1/2$, this equals the mean absolute error of the median forecast.

Value

A matrix of the quantile scores for the desired quantiles, one column per quantile.

References

Gneiting, T. and Raftery, A. E. (2007), 'Strictly proper scoring rules, prediction, and estimation', Journal of the American Statistical Association 102(477), 359-378

See Also

predict.idrcal, qpred

Examples

```
data("rain")

## IDRsd based on ensemble forecast

ensemble <- rain[1:(3 * 365), 3:54, drop = FALSE]
y <- rain[1:(3 * 365), "obs"]

fit <- idrsd(y = y, X = ensemble, type = 'ensemble')

## Compute mean absolute error of the median postprocessed forecast using
## data of the next year (out-of-sample predictions) and compare to raw
## forecast

ytest = rain[(3 * 365 + 1):(4 * 365), "obs"]
ensemble_test = rain[(3 * 365 + 1):(4 * 365), 3:54,drop = FALSE]
predictions <- predict(fit, data = ensemble_test)

idrMAE <- mean(qscore(predictions, 0.5, ytest))
rawMAE <- mean(qscore(ensemble_test, 0.5, ytest))

c("idrsd" = idrMAE, "raw" = rawMAE)
```

Index